

2. Post_Clean_NB

July 25, 2018

1 Naive Bayes on Amazon Reviews Dataset (Part II)

1.1 Data Source:

The preprocessing step has produced `final.sqlite` file after doing the data preparation & cleaning. The review text is now devoid of punctuations, HTML markups and stop words.

1.2 Objective:

To find Precision, Recall, F1 Score, Confusion Matrix, Accuracy of 10-fold cross validation Naive Bayes on vectorized input data, for BoW and TF-IDF featurizations. TPR, TNR, FPR and FNR is calculated for both. The most frequent words in positive and negative class also needs to be found out

1.3 At a glance:

Random Sampling is done to reduce input data size and time based slicing to split into training and testing data. **The metrics obtained by applying 10-fold cross validation with Naive Bayes using BoW and tf-idf featurizations are compared to find optimal alpha for smoothing.**

The accuracy for different alpha values are plotted. The performance metrics with both featurizations are computed. **Most frequent words** in both classes are also enumerated.

1.4 Laplace Smoothing:

Range of Alpha: taken values ranging from 10^{-6} to 10^3

To try out different values of alpha, **Geometric Progression is used as step function** in the above alpha range.

2 Preprocessed Data Loading

```
In [42]: #loading libraries for NB
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
# from sklearn.metrics import accuracy_score
```

```

# import sklearn.metrics
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation

#loading libraries for scikit learn, nlp, db, plot and matrix.
import sqlite3
import pdb
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

# using the SQLite Table to read data.
con = sqlite3.connect('./final.sqlite')

#filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
final = pd.read_sql_query("""
SELECT *
FROM Reviews
""", con)

print(final.head(3))
print(final.shape)

```

	index	Id	ProductId	UserId	ProfileName	\
0	138706	150524	0006641040	ACITT7DI6IDDL	shari zychinski	
1	138688	150506	0006641040	A2IW4PEEK02R0U	Tracy	
2	138689	150507	0006641040	A1S4A3IQ2MU7V4	sally sue	"sally sue"

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	0	0	positive	939340800	
1	1	1	positive	1194739200	
2	1	1	positive	1191456000	

```

                                Summary \
0          EVERY book is educational
1 Love the book, miss the hard cover version
2          chicken soup with rice months

                                Text \
0 this witty little book makes my son laugh at l...
1 I grew up reading these Sendak books, and watc...
2 This is a fun way for children to learn their ...

                                CleanedText
0 b'witti littl book make son laugh loud recit c...
1 b'grew read sendak book watch realli rosi movi...
2 b'fun way children learn month year learn poem...
(364171, 12)

```

3 Random Sampling & Time Based Slicing

```

In [43]: # To randomly sample the data and sort based on time before doing train/ test split.
         # The slicing into train & test data will be done later in kfoldcv() function.

num_points = 200000

# used to format headings
bold = '\033[1m'
end = '\033[0m'

# you can use random_state for reproducibility
sampled_final = final.sample(n=num_points, random_state=2)

#Sorting data according to Time in ascending order
sorted_final = sampled_final.sort_values('Time', axis=0,
                                         ascending=True, inplace=False, kind='quicksort', na_position='last')

# fetching the outcome class
y = sorted_final['Score'] # showing you two ways of indexing a pandas df
print(y.shape)

X_train, X_test, y_train, y_test = train_test_split(
    sorted_final, y, test_size=0.3, random_state=42)

(200000,)

```

4 Custom Defined Functions

2 user defined functions are written to

- a) K-fold Cross Validation & estimation of Optimal Alpha.
- b) Compute Performance Metrics of NB Classifier.
- c) Find Most Frequent Words.

4.1 a) k-fold Cross Validation & Optimal Alpha Estimation

```
In [44]: # split the data set into train and test. Do 10-fold cross validation
# For Binary BoW representation, Bernoulli NB is used.
# For count based BoW and tf-idf, Multinomial NB is used.
```

```
import numpy
import math
import matplotlib.pyplot as plt
from sklearn.naive_bayes import BernoulliNB, MultinomialNB

def kfoldcv(X_train_vect, X_test_vect,
            algo='MultinomialNB', toPlot = False, title_cf=''):

    # Time based slicing of data into train and test.
    #     num_train_data = int(split_ratio_train*X.shape[0])

    #     X_train = X[0:num_train_data]
    #     y_train = y[0:num_train_data]
    #     X_test = X[num_train_data+1:]
    #     y_test = y[num_train_data+1:]

    # generate gp sequence =  $a \cdot r^n$ 
    a = 10**-6
    r = 2
    n = int(math.log(10**3/a, r)) # of times to do gp to reach  $10^3$ 
    alphas = [a * r**i for i in range(n)]

    # empty list that will hold cv scores
    cv_scores = []

    # perform 10-fold cross validation
    for a in alphas:
        if (algo=='MultinomialNB'):
            nb = MultinomialNB(alpha=a)
        else:
```

```

        nb = BernoulliNB(alpha=a)

    scores = cross_val_score(nb, X_train_vect, y_train, cv=10)
    cv_scores.append(scores.mean())

    if (toPlot):
        plt.figure()
        plt.plot(alphas, cv_scores)
        plt.xlabel('Alpha Values')
        plt.ylabel('Accuracy Obtained')
        plt.title('Cross Validation Plot: Alpha vs Accuracy')

    # changing to misclassification error
    MSE = [1 - x for x in cv_scores]

    # determining best k
    optimal_alpha = alphas[MSE.index(min(MSE))]
    print('\nThe optimal value of alpha is %f.' % optimal_alpha)

    return compute_accuracy(X_train_vect, X_test_vect,
                           alpha_val=optimal_alpha, algo=algo, title_cf=title_cf)

```

4.2 b) Compute NB Classifier Performance Metrics

```

In [45]: # ===== NB with alpha = optimal_alpha =====
#To compute the performance metrics of NB classifier
# For Binary BoW representation, Bernoulli NB is used.
# For count based BoW and tf-idf, Multinomial NB is used.

import seaborn as sn
from sklearn.metrics import *

def compute_accuracy(X_train_vect, X_test_vect,
                    alpha_val, algo = 'BernoulliNB', title_cf = 'Confusion Matrix'):

    # Time based slicing of data into train and test.
    #     num_train_data = int(split_ratio_train*X.shape[0])

    #     X_train = X[0:num_train_data]
    #     y_train = y[0:num_train_data]
    #     X_test = X[num_train_data+1:]
    #     y_test = y[num_train_data+1:]

    # instantiate learning model k = optimal_k
    if (algo == 'BernoulliNB'):
        nb_optimal = BernoulliNB(alpha=alpha_val)
    else:
        nb_optimal = MultinomialNB(alpha=alpha_val)

```

```

# fitting the model
nb_optimal.fit(X_train_vect, y_train)

# predict the response
pred = nb_optimal.predict(X_test_vect)

print(bold + '\n\nMetric Analysis of NB Classifier for Optimal Alpha' + end)
# evaluate accuracy
acc = accuracy_score(y_test, pred) * 100
print('\nAccuracy = %f' % acc)

precision = precision_score(y_test, pred, pos_label='positive') * 100
print('\nPrecision = %f' % precision)

recall = recall_score(y_test, pred, pos_label='positive') * 100
print('\nRecall = %f' % recall)

f1score = f1_score(y_test, pred, pos_label='positive') * 100
print('\nF1 Score = %f' % f1score)

confusion = confusion_matrix(y_test, pred, labels=["positive", "negative"])
print(bold + "\n\nConfusion Matrix" + end)
print(confusion)

plt.figure()
plt.title(title_cf)
df_cm = pd.DataFrame(confusion, range(2), range(2))
sn.set(font_scale=1.4)#for label size
sn.heatmap(df_cm, annot=True,annot_kws={"size": 16}, fmt="d")# font size

(tn, fp, fn, tp) = confusion.ravel()
print("\nTrue Negatives = " + str(tn))
print("True Positives = " + str(tp))
print("False Negatives = " + str(fn))
print("False Positives = " + str(fp))

actual_positives = tp+fn
actual_negatives = tn+fp
print("\nTotal Actual Positives = " + str(actual_positives))
print("Total Actual Negatives = " + str(actual_negatives))

print("\nTrue Positive Rate(TPR) = " + str(round(tp/actual_positives, 2)))
print("True Negative Rate(TNR) = " + str(round(tn/actual_negatives, 2)))
print("False Positive Rate(FPR) = " + str(round(fp/actual_negatives, 2)))
print("False Negative Rate(FNR) = " + str(round(fn/actual_positives, 2)))

```

```
return nb_optimal.feature_log_prob_
```

4.3 c) Find Most Frequent Words

In [46]: *# To find out the out positive and negative words based on feature log probabilities*

```
# To get important words from feature_log_prob
# 1)Sort log prob values in ascending order and get the indices
# 2)Get the ending 'n' indices to find top n feature names (nwords)
# having most occurrences in class 1 and class 0
# 3)Use BoW.get_feature_name and find words corresponding to above 'n' indices

def find_top_words(vect, feature_log_probs, nwords):

    neg_class_prob_sorted = feature_log_probs[0, :].argsort()
    pos_class_prob_sorted = feature_log_probs[1, :].argsort()

    top_neg_words = np.take(vect.get_feature_names(),
                            neg_class_prob_sorted[neg_class_prob_sorted.size-nwords:])
    top_pos_words = np.take(vect.get_feature_names(),
                            pos_class_prob_sorted[pos_class_prob_sorted.size-nwords:])

    print(bold + "\n\nTop Negative Words: "+ end)
    for id, word in enumerate(top_neg_words):
        print("\t" + word + "\t Log Prob: " + str(
            round(feature_log_probs[
                0, neg_class_prob_sorted[neg_class_prob_sorted.size-nwords+id]], 2)))

    print(bold + "\n\nTop Positive Words: "+ end)
    for id, word in enumerate(top_pos_words):
        print("\t" + word + "\t Log Prob: " + str(
            round(feature_log_probs[
                1, pos_class_prob_sorted[pos_class_prob_sorted.size-nwords+id]], 2)))
```

5 Bernoulli/ Multinomial NB on Binary BoW/ Count BoW

BoW will result in a **sparse matrix with huge number of features** as it creates a feature for each unique word in the review.

For Binary BoW feature representation, Bernoulli NB is used as the value can take only 0 and 1. For count based BoW and tf-idf, Multinomial NB is used. The variation of accuracy corresponding to varying values of alpha is plotted and the alpha with the highest accuracy is identified. Top words in both classes are found out using log probabilities.

In [47]: *#BoW*

```

# from sklearn.decomposition import TruncatedSVD
from sklearn.random_projection import sparse_random_matrix

# X_train, X_test, y_train, y_test

#Binary BoW
count_vect_bin = CountVectorizer(binary = True) #in scikit-learn
X_bin_train_vect = count_vect_bin.fit_transform(X_train['CleanedText'].values)
X_bin_train_vect.get_shape()

#BoW
count_vect = CountVectorizer() #in scikit-learn
X_train_vect = count_vect.fit_transform(X_train['CleanedText'].values)
X_train_vect.get_shape()

#BoW Test
# count_vect = CountVectorizer() #in scikit-learn
X_test_bin_vect = count_vect_bin.transform(X_test['CleanedText'].values)
# X_test_bin_vect.get_shape()

#BoW Test
# count_vect = CountVectorizer() #in scikit-learn
X_test_vect = count_vect.transform(X_test['CleanedText'].values)
# X_test_vect.get_shape()

# print(X_test_bin_vect.get_shape())
print(X_train_vect.get_shape())

print(bold + "\n\n1) BoW with Multinomial NB" + end)
feature_logprobs = kfoldcv(X_train_vect, X_test_vect,
    algo = 'MultinomialNB', toPlot=True,
    title_cf='Count BoW + Multinomial NB Confusion Matrix HeatMap')

# To run Binary BoW with Bernoulli NB
print(bold + "\n\n2) Binary BoW with Bernoulli NB" + end)
kfoldcv(X_bin_train_vect, X_test_bin_vect, algo="BernoulliNB",
    title_cf='Binary BoW + Bernoulli NB Confusion Matrix HeatMap')

# To print 50 top words - positive and negative
find_top_words(count_vect, feature_logprobs, 50)

```

(140000, 44920)

1) BoW with Multinomial NB

The optimal value of alpha is 2.097152.

Metric Analysis of NB Classifier for Optimal Alpha

Accuracy = 90.620000

Precision = 92.661611

Recall = 96.530137

F1 Score = 94.556323

Confusion Matrix

[[48879 1757]

[3871 5493]]

True Negatives = 48879

True Positives = 5493

False Negatives = 3871

False Positives = 1757

Total Actual Positives = 9364

Total Actual Negatives = 50636

True Positive Rate(TPR) = 0.59

True Negative Rate(TNR) = 0.97

False Positive Rate(FPR) = 0.03

False Negative Rate(FNR) = 0.41

2) Binary BoW with Bernoulli NB

The optimal value of alpha is 0.008192.

Metric Analysis of NB Classifier for Optimal Alpha

Accuracy = 89.280000

Precision = 93.122488

Recall = 94.259025

F1 Score = 93.687310

Confusion Matrix

[[47729 2907]

[3525 5839]]

True Negatives = 47729

True Positives = 5839

False Negatives = 3525

False Positives = 2907

Total Actual Positives = 9364

Total Actual Negatives = 50636

True Positive Rate(TPR) = 0.62

True Negative Rate(TNR) = 0.94

False Positive Rate(FPR) = 0.06

False Negative Rate(FNR) = 0.38

Top Negative Words:


sugar	Log Prob: -6.05
give	Log Prob: -6.05
know	Log Prob: -5.99
didnt	Log Prob: -5.99
drink	Log Prob: -5.98
made	Log Prob: -5.97
say	Log Prob: -5.96
water	Log Prob: -5.95
could	Log Prob: -5.94
price	Log Prob: -5.93
think	Log Prob: -5.93
chocol	Log Prob: -5.93
also	Log Prob: -5.92
better	Log Prob: -5.91
bad	Log Prob: -5.89
want	Log Prob: -5.86
bought	Log Prob: -5.84
first	Log Prob: -5.84
disappoint	Log Prob: -5.83
purchas	Log Prob: -5.76
review	Log Prob: -5.75
packag	Log Prob: -5.72
love	Log Prob: -5.71
look	Log Prob: -5.69
dog	Log Prob: -5.67
bag	Log Prob: -5.64
realli	Log Prob: -5.63
much	Log Prob: -5.61
eat	Log Prob: -5.59
make	Log Prob: -5.58
time	Log Prob: -5.58
amazon	Log Prob: -5.56
box	Log Prob: -5.49
even	Log Prob: -5.45
tea	Log Prob: -5.43
dont	Log Prob: -5.38
food	Log Prob: -5.34
order	Log Prob: -5.32
buy	Log Prob: -5.24
get	Log Prob: -5.24
coffe	Log Prob: -5.17
good	Log Prob: -5.14
use	Log Prob: -5.12
tri	Log Prob: -4.99
would	Log Prob: -4.97

flavor	Log Prob: -4.88
one	Log Prob: -4.82
product	Log Prob: -4.53
like	Log Prob: -4.4
tast	Log Prob: -4.32


Top Positive Words:

high	Log Prob: -5.97
first	Log Prob: -5.97
cup	Log Prob: -5.97
found	Log Prob: -5.94
sugar	Log Prob: -5.91
box	Log Prob: -5.91
water	Log Prob: -5.89
sweet	Log Prob: -5.88
recommend	Log Prob: -5.85
year	Log Prob: -5.85
chocol	Log Prob: -5.8
day	Log Prob: -5.8
better	Log Prob: -5.79
mix	Log Prob: -5.78
even	Log Prob: -5.72
ive	Log Prob: -5.68
bag	Log Prob: -5.67
store	Log Prob: -5.65
dog	Log Prob: -5.64
drink	Log Prob: -5.63
well	Log Prob: -5.61
littl	Log Prob: -5.56
order	Log Prob: -5.54
dont	Log Prob: -5.53
best	Log Prob: -5.52
find	Log Prob: -5.51
also	Log Prob: -5.51
much	Log Prob: -5.51
price	Log Prob: -5.49
amazon	Log Prob: -5.48
realli	Log Prob: -5.44
eat	Log Prob: -5.43
buy	Log Prob: -5.38
time	Log Prob: -5.37
would	Log Prob: -5.36
food	Log Prob: -5.2
get	Log Prob: -5.1
make	Log Prob: -5.07
coffe	Log Prob: -5.02
tea	Log Prob: -4.96
tri	Log Prob: -4.92
product	Log Prob: -4.87


one	Log Prob: -4.81
great	Log Prob: -4.74
use	Log Prob: -4.74
love	Log Prob: -4.7
flavor	Log Prob: -4.69
good	Log Prob: -4.66
tast	Log Prob: -4.52
like	Log Prob: -4.45



Count BoW + Multinomial NB Confusion Matrix HeatMap



Binary BoW + Bernoulli NB Confusion Matrix HeatMap



6 Multinomial NB on tf-IDF Featurization

Sparse matrix generated from tf-IDF is fed in to Multinomial Naive Bayes to find the optimal alpha value. Performance metrics of Multinomial NB with tf-idf featurization is found.

```
In [48]: #TF-IDF
tf_idf_vect = TfidfVectorizer()
X_train_vect = tf_idf_vect.fit_transform(X_train['CleanedText'].values)
X_train_vect.get_shape()

#TF-IDF Test
# tf_idf_vect = TfidfVectorizer() #in scikit-learn
X_test_vect = tf_idf_vect.transform(X_test['CleanedText'].values)
X_test_vect.get_shape()

print(X_train_vect.get_shape())

print(bold + "\n TF-IDF with Multinomial NB" + end)
# To run brute & kd-tree knn & also time the code
feature_logprobs = kfoldcv(X_train_vect, X_test_vect, algo = 'MultinomialNB',
                           title_cf='Count BoW + Multinomial NB Confusion Matrix HeatMap')

# To print 50 top words - positive and negative
find_top_words(tf_idf_vect, feature_logprobs, 50)

(140000, 44920)
TF-IDF with Multinomial NB
```

The optimal value of alpha is 0.016384.

Metric Analysis of NB Classifier for Optimal Alpha

Accuracy = 87.821667

Precision = 87.863747

Recall = 99.283119

F1 Score = 93.225038

Confusion Matrix

```
[[50273  363]
 [ 6944 2420]]
```

True Negatives = 50273

True Positives = 2420

False Negatives = 6944

False Positives = 363

Total Actual Positives = 9364

Total Actual Negatives = 50636

True Positive Rate(TPR) = 0.26

True Negative Rate(TNR) = 0.99

False Positive Rate(FPR) = 0.01

False Negative Rate(FNR) = 0.74

Top Negative Words:

say	Log Prob: -6.18
wast	Log Prob: -6.17
water	Log Prob: -6.16
first	Log Prob: -6.14
got	Log Prob: -6.14
think	Log Prob: -6.14
better	Log Prob: -6.14
price	Log Prob: -6.1
item	Log Prob: -6.1
smell	Log Prob: -6.1
could	Log Prob: -6.1
make	Log Prob: -6.09
want	Log Prob: -6.09
thought	Log Prob: -6.05
receiv	Log Prob: -6.02
chocol	Log Prob: -6.0
didnt	Log Prob: -6.0
time	Log Prob: -5.95
review	Log Prob: -5.94
money	Log Prob: -5.92
look	Log Prob: -5.92
much	Log Prob: -5.91
realli	Log Prob: -5.91
amazon	Log Prob: -5.91
bought	Log Prob: -5.91
packag	Log Prob: -5.88
eat	Log Prob: -5.88
dog	Log Prob: -5.88
purchas	Log Prob: -5.85
bag	Log Prob: -5.84
bad	Log Prob: -5.83
food	Log Prob: -5.79
even	Log Prob: -5.77
use	Log Prob: -5.71
get	Log Prob: -5.67
dont	Log Prob: -5.67
tea	Log Prob: -5.66
good	Log Prob: -5.66
box	Log Prob: -5.62
disappoint	Log Prob: -5.61
order	Log Prob: -5.52


buy	Log Prob: -5.49
tri	Log Prob: -5.49
coffe	Log Prob: -5.37
one	Log Prob: -5.37
flavor	Log Prob: -5.35
would	Log Prob: -5.32
like	Log Prob: -5.01
product	Log Prob: -5.01
tast	Log Prob: -4.84

Top Positive Words:

sugar	Log Prob: -6.15
snack	Log Prob: -6.15
nice	Log Prob: -6.15
favorit	Log Prob: -6.13
enjoy	Log Prob: -6.12
even	Log Prob: -6.12
year	Log Prob: -6.11
day	Log Prob: -6.1
treat	Log Prob: -6.07
sweet	Log Prob: -6.07
mix	Log Prob: -6.05
better	Log Prob: -6.04
delici	Log Prob: -6.02
recommend	Log Prob: -6.02
ive	Log Prob: -5.99
dont	Log Prob: -5.98
bag	Log Prob: -5.97
also	Log Prob: -5.96
well	Log Prob: -5.96
chocol	Log Prob: -5.96
littl	Log Prob: -5.9
drink	Log Prob: -5.9
much	Log Prob: -5.88
store	Log Prob: -5.83
would	Log Prob: -5.81
dog	Log Prob: -5.79
eat	Log Prob: -5.79
realli	Log Prob: -5.75
amazon	Log Prob: -5.75
time	Log Prob: -5.74
find	Log Prob: -5.74
order	Log Prob: -5.72
food	Log Prob: -5.68
best	Log Prob: -5.67
buy	Log Prob: -5.66
price	Log Prob: -5.64
get	Log Prob: -5.61
make	Log Prob: -5.56

tri	Log Prob: -5.47
one	Log Prob: -5.41
use	Log Prob: -5.32
product	Log Prob: -5.3
coffe	Log Prob: -5.22
flavor	Log Prob: -5.22
tea	Log Prob: -5.21
like	Log Prob: -5.15
good	Log Prob: -5.14
tast	Log Prob: -5.13
love	Log Prob: -5.07
great	Log Prob: -5.06

Count BoW + Multinomial NB Confusion Matrix HeatMap



```
In [50]: from IPython.display import Image
         Image(filename='summary.jpg')
```

Out[50]:

Performance Metric Summary

Model	Best Hyper Parameter	Train metric	Test metric (A = Accuracy; P = Precision; R = Recall; F1 = F1 Score; TPR = True +ve Rate; TNR = True -ve Rate; FPR = False +ve; FNR = False -ve)
Binary BoW on Bernoulli NB	0.008192	70000 reviews; 44920 features	A = 89.28; P = 93.12; R = 94.26; F1 = 93.68 ; TPR = 0.62; TNR = 0.94; FPR = 0.06; FNR = 0.38
Multinomial NB on Count based BoW	2.097152	70000 reviews; 44920 features	A = 90.62; P = 92.66; R = 96.53; F1 = 94.56 ; TPR = 0.59; TNR = 0.97; FPR = 0.03; FNR = 0.41
Multinomial Naive Bayes on tf-IDF	0.016384	70000 reviews; 44920 features	A = 87.82; P = 87.86; R = 99.28; F1 = 93.23 ; TPR = 0.26; TNR = 0.99; FPR = 0.01; FNR = 0.74

6.1 Observations

- 1) F1 Score of **"Count based BoW with Multinomial NB"** is slightly higher than **"Binary BoW with Bernoulli NB"** method. This indicates the **loss of information** when the count vector is made binary.
- 2) Accuracy of TF-IDF method is 2% less than that we got from BoW method.
- 3) TNR and Recall of **TF-IDF with Multinomial NB** method is as high as 99%. This indicates the **negative reviews are identified properly** with only 1% false positive rate.
- 4) If we need a system with a **prime requirement to correctly identify as many negative reviews** as possible, then **"Multinomial NB on tf-IDF"** should be used (as TNR = 0.99)
- 5) TPR of **TF-IDF with Multinomial NB** method is low (26%). 74% of **positive reviews are not identified** correctly (FPR = 0.74)
- 6) TPR of **"BoW with Multinomial NB"** has the highest F1 Score, amongst all the three methods. Hence, **Bag of Words** featurization with multinomial Naive Bayes is the classifier of choice.
- 7) Naive Bayes algorithm is based on Bayes' theorem with a strong (naive) conditional independence assumption between the features. Hence, **working with w2v features which are completely dependent is not a good idea**. Naive Bayes on W2V and tf-idf weighted W2v is not done for this reason.